

Introduction to OpenTelemetry with Flask

FlaskCon 2024



Jessica Garson

@JessicaGarson

@JessicaGarson@macaw.social

Senior Developer Advocate



What is Elastic?

<https://github.com/JessicaGarson/Introduction-to-OpenTelemetry-with-Flask>



I had attempted to find a short primer on starting with OpenTelemetry with Flask but couldn't find much more than a [demo application](#).



ChatGPT

Sure! Imagine you have a big box of crayons (these are like different parts of a computer program) and you want to make a beautiful picture (which is like running the program to do something cool). But sometimes, while you're drawing, a crayon might break or not color the way it should. This can make your picture not turn out as nice as you wanted.

OpenTelemetry is like having a magical friend who watches you while you draw. This friend takes notes on which crayons you use, how they perform, and if they break. Then, if something goes wrong, your magical friend can tell you exactly which crayon broke and why, so you can fix it and make sure your picture turns out great next time!

So, OpenTelemetry helps people who make computer programs make sure their programs run smoothly and fix any problems that happen.



Observability fundamentals

Observability helps you determine what's happening inside your system by analyzing the data it outputs



What isn't operating as expected

Beyond detecting problems

What is an agent?

An observability agent is a software component or tool that collects, processes, and sends telemetry data from a system or application to an observability platform for analysis.

What is OpenTelemetry (OTel)?

OpenTelemetry (OTel) is an observability framework that is open source and vendor neutral, designed to work with any backend system.

It provides standardized APIs, libraries, and tools to collect telemetry data such as metrics, logs, and traces.

OpenTelemetry enables comprehensive monitoring
across services and platforms.

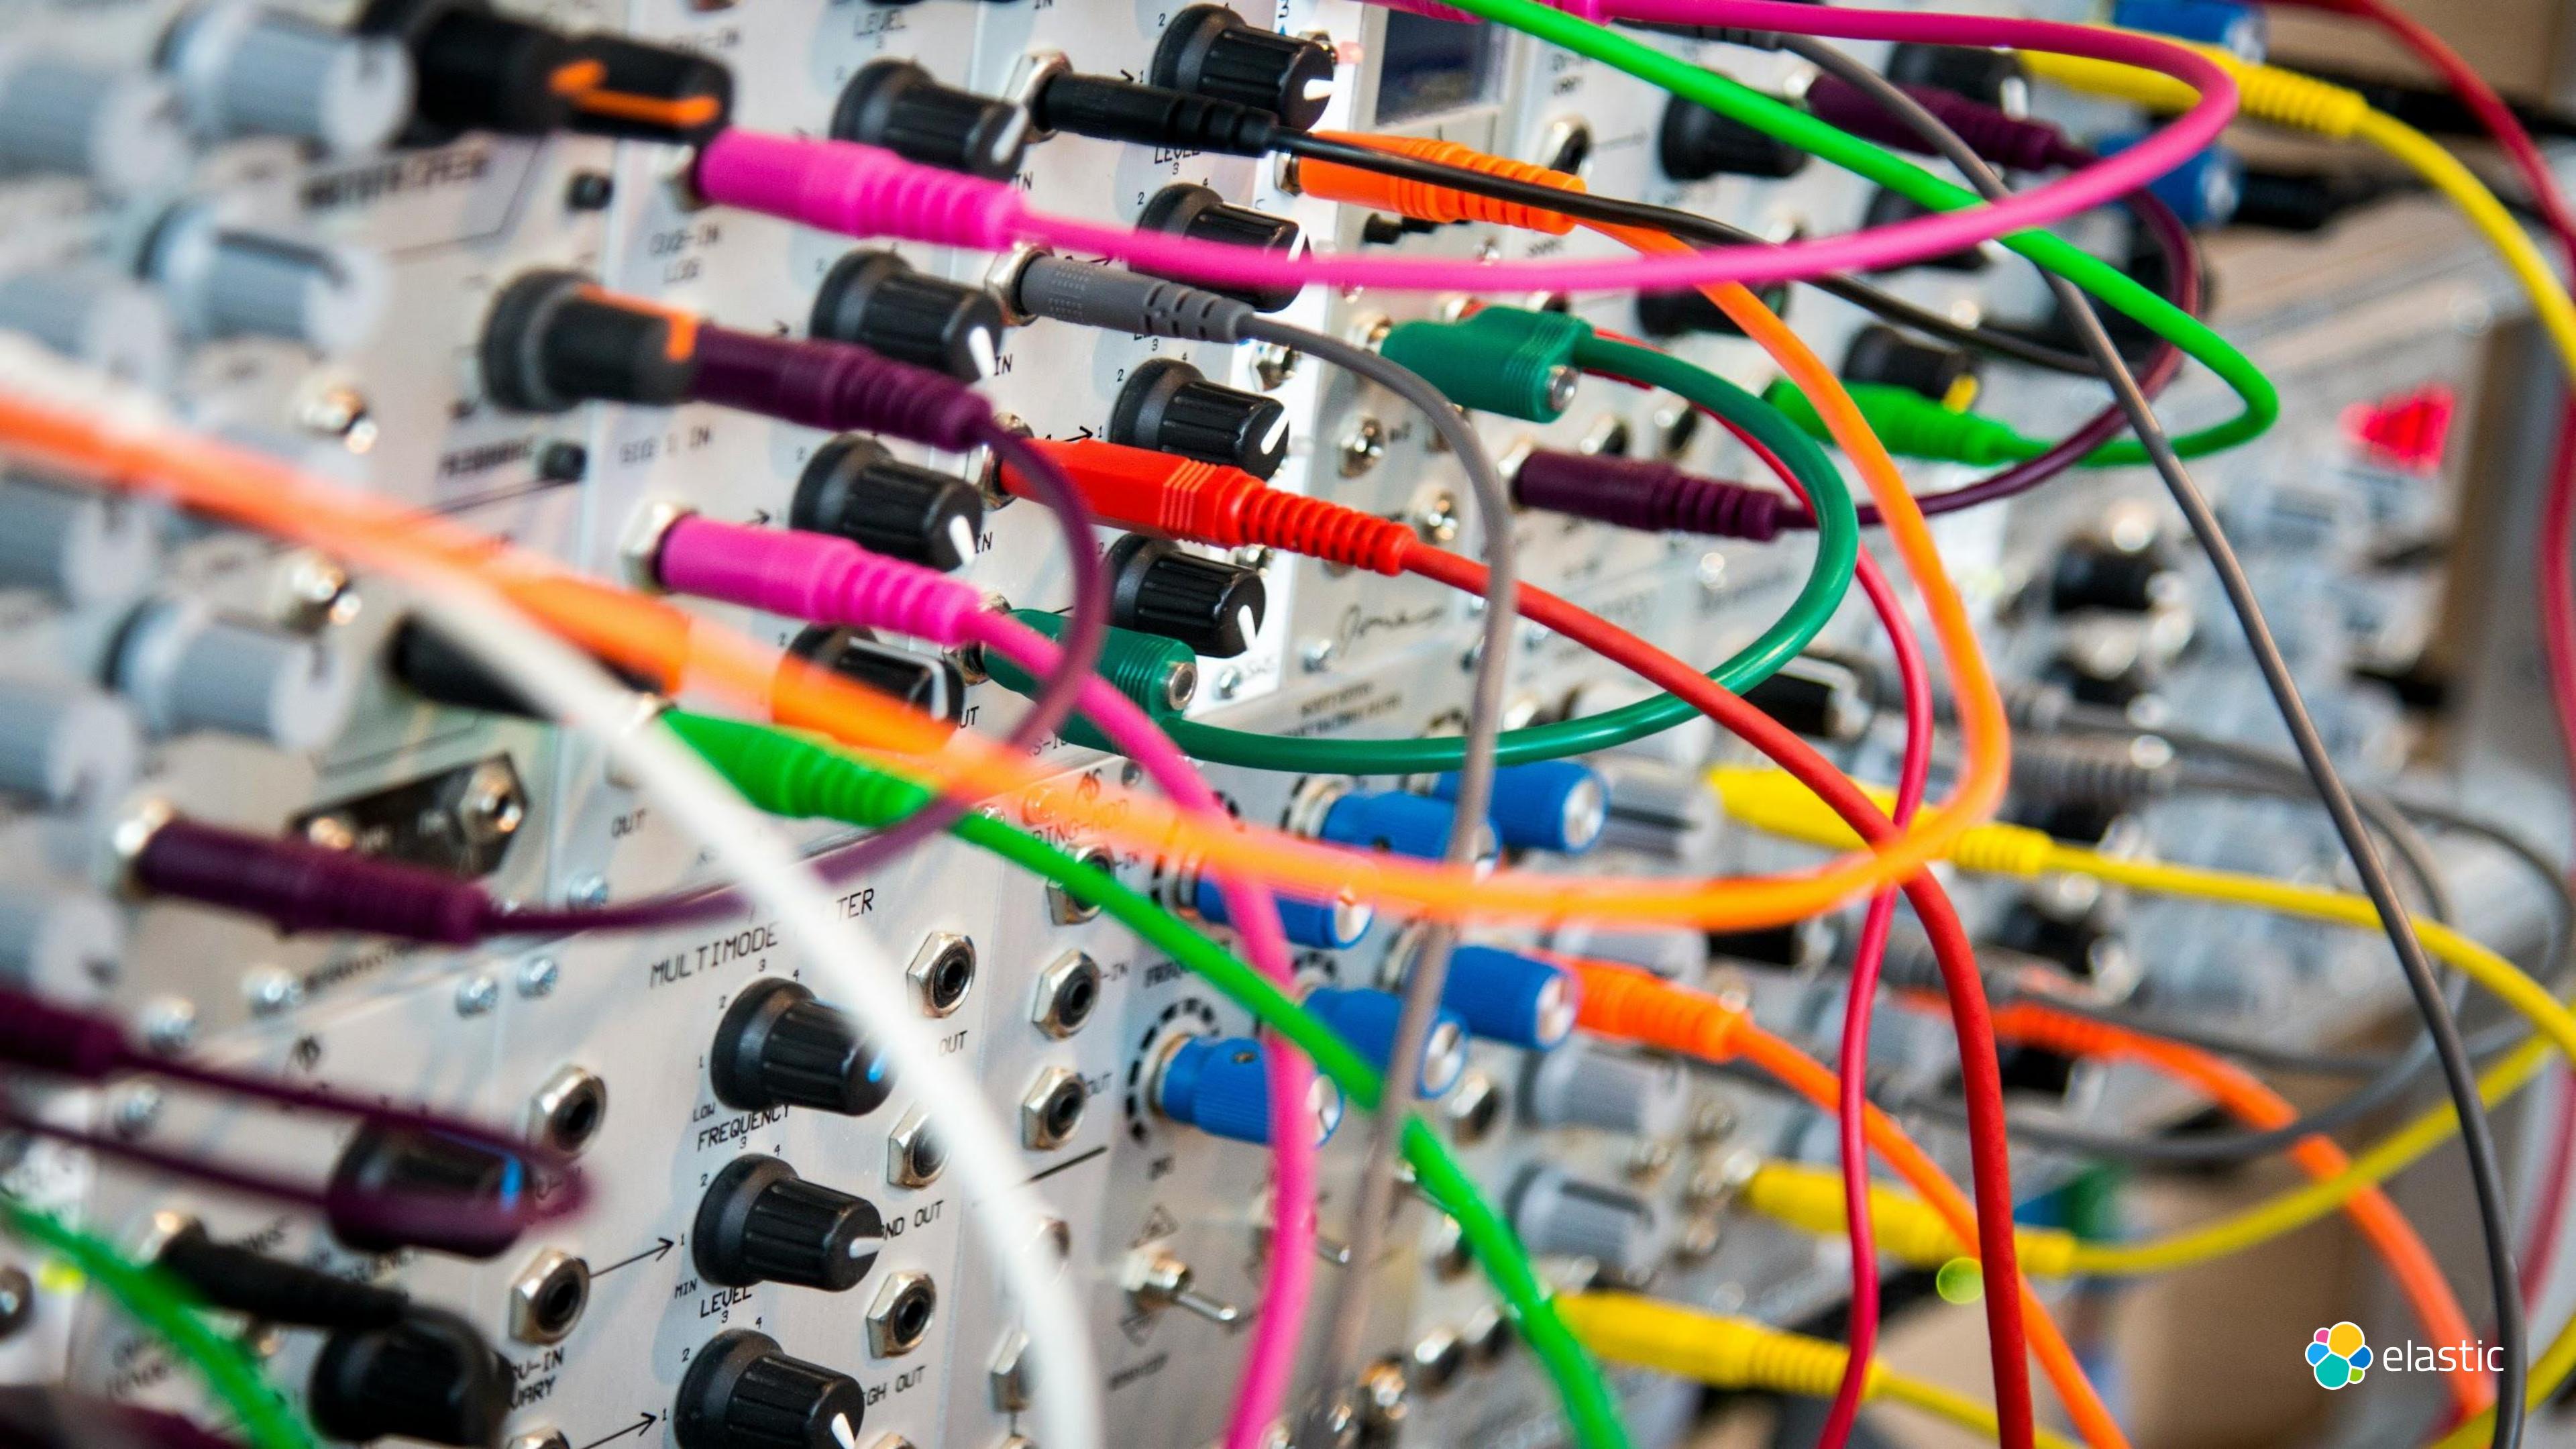
Why?

A photograph of a stack of logs in a forest. The logs are cut trees, showing their circular cross-sections. They are piled up in a somewhat haphazard manner, with some logs leaning against each other. The background is filled with dense forest trees, and the ground is covered with fallen branches and needles. The lighting suggests it might be late afternoon or early morning.

Logs alone are simply not
enough sometimes

Um, what?

| | | | | | | | |
|---|--|------------|--------------|------------------------|----------|-----|---|
| > | * | 2024-05-10 | 08:00:06.117 | 55 | 54440319 | ... | http://api.nasa.gov/neo/rest/v1/neo/sentry/544... |
| > | * | 2024-05-10 | 08:00:06.117 | 56 | 54440642 | ... | NaN |
| ▼ | * | 2024-05-10 | 08:00:06.117 | 57 | 54440644 | ... | http://api.nasa.gov/neo/rest/v1/neo/sentry/544... |
| | | | | | | | |
| ▼ | { | | | | | | |
| | insertId: "663e0c460001c9a60e5e65c3" | | | | | | |
| ▶ | labels: {2} | | | | | | |
| | logName: "projects/ncav-293119/logs/run.googleapis.com%2Fstdout" | | | | | | |
| | receiveTimestamp: "2024-05-10T12:00:06.338017672Z" | | | | | | |
| ▶ | resource: {2} | | | | | | |
| | textPayload: "57 54440644 ... http://api.nasa.gov/neo/rest/v1/neo/sentry/544..." | | | | | | |
| | timestamp: "2024-05-10T12:00:06.117158Z" | | | | | | |
| | } | | | | | | |
| > | * | 2024-05-10 | 08:00:06.117 | [43 rows x 18 columns] | | | |



Future proof



Benefits

- Enhanced visibility into the execution flow and call timing
- Metrics on how users utilize services to pinpoint unnecessary ones
- The ability to detect patterns, such as peak usage times, which aids in production planning and scaling.

Logs

Logs are records of events in a system, documenting operations, errors, and activities to aid in troubleshooting, monitoring, and compliance.

Metrics

Metrics are quantitative measurements that track the performance and health of a system

Traces

Traces track the path and interactions of a request through a system.

Spans within traces? What?!?!



Spans are the building blocks of traces.
They include the following information:

- Name
- Parent span ID (empty for root spans)

Ohhhh! A span within a trace!

Span #1

```
Trace ID      : 4e6b8d4dcd44d9245285a9459a0114e6
Parent ID     : 6c2171f87b96401f
ID            : d56a8dc2bb6f2dbb
Name          : jinja2.render
Kind          : Internal
Start time    : 2024-05-14 19:07:24.897193 +0000 UTC
End time      : 2024-05-14 19:07:24.897258 +0000 UTC
Status code   : Unset
Status message :
```

Attributes:

```
-> jinja2.template_name: Str(<memory>)
    {"kind": "exporter", "data_type": "traces", "name": "debug"}
```

Root span

Span #1

| | | |
|----------------|---|--------------------------------------|
| Trace ID | : | 4e6b8d4dcd44d9245285a9459a0114e6 |
| Parent ID | : | |
| ID | : | 6c2171f87b96401f |
| Name | : | / |
| Kind | : | Server |
| Start time | : | 2024-05-14 19:07:24.894406 +0000 UTC |
| End time | : | 2024-05-14 19:07:24.897374 +0000 UTC |
| Status code | : | Unset |
| Status message | : | |

The classic Flask example

To-Do List

ADD TASK

Cool! A new task!

Delete

Let's instrument our app

Instrumenting refers to the process of adding observability features to your application to collect telemetry data, such as traces, metrics, and logs.

The command line is just one of the ways you can instrument your application. Other options include serverless functions and kubernetes.



```
pip install opentelemetry-distro  
opentelemetry-bootstrap -a install
```



```
export OTEL_PYTHON_LOGGING_AUTO_INSTRUMENTATION_ENABLED=true
opentelemetry-instrument \
    --traces_exporter console \
    --metrics_exporter console \
    --logs_exporter console \
    --service_name todo \
flask run -p 8080
```

So when would you want to use manual instrumentation?

Logs



import logging



```
# Initialize the logger
logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(levelname)s - %(message)s'
)
```



```
@app.route('/', methods=['GET'])
def home():
    tasks = Task.query.all()
    logging.info(f'Fetched {len(tasks)} tasks from the database')
    return render_template_string(HOME_HTML, tasks=tasks)

@app.route('/add', methods=['POST'])
def add():
    task_description = request.form['task']
    new_task = Task(description=task_description)
    db.session.add(new_task)
    db.session.commit()
    logging.info(f'Added new task with description: {task_description}')
    return redirect(url_for('home'))
```



Traces



```
from opentelemetry import trace
```



```
tracer = trace.get_tracer("task.tracer")
```



```
@app.route('/', methods=['GET'])
def home():
    with tracer.start_as_current_span("querying_tasks"):
        tasks = Task.query.all()
    return render_template_string(HOME_HTML, tasks=tasks)

@app.route('/add', methods=['POST'])
def add():
    task_description = request.form['task']
    new_task = Task(description=task_description)
    with tracer.start_as_current_span("adding_task"):
        db.session.add(new_task)
        db.session.commit()
    return redirect(url_for('home'))
```

Metrics



```
from opentelemetry import metrics
```



```
meter = metrics.get_meter("task.meter")

request_counter = meter.create_counter(
    "flask_app_requests_total",
    description="Total number of requests",
)

task_counter = meter.create_counter(
    "flask_app_tasks_total",
    description="Total number of tasks created",
)
```



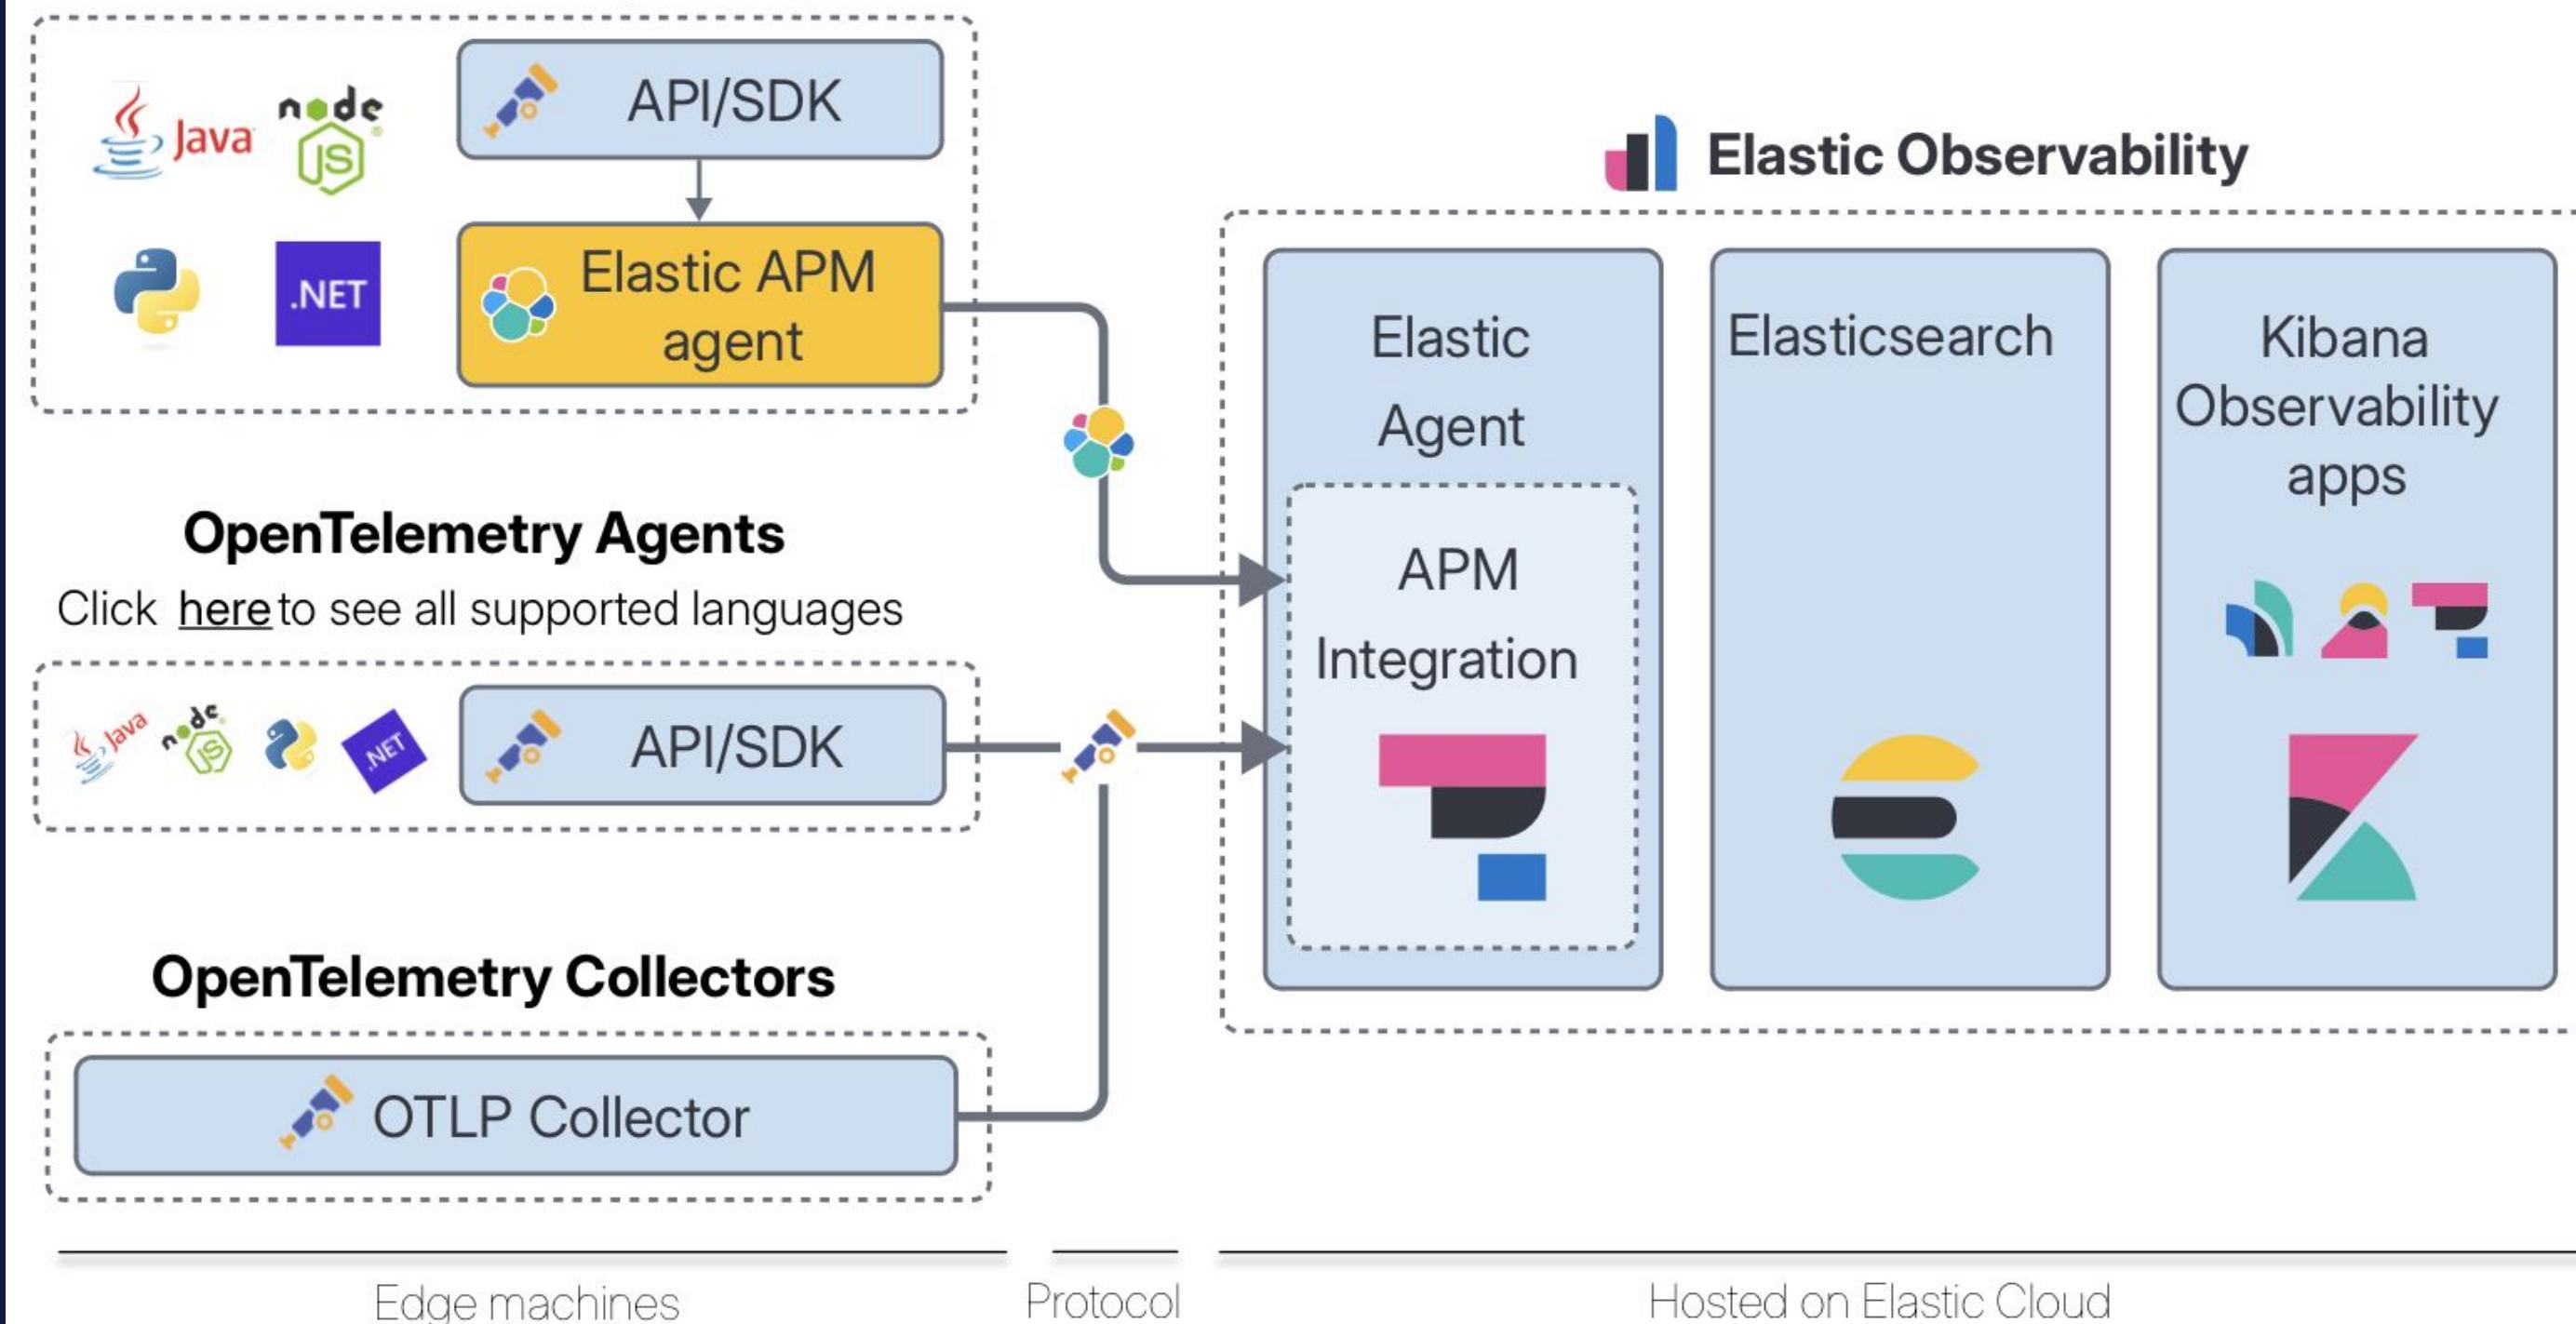
```
@app.route('/add', methods=['POST'])
def add():
    task_description = request.form['task']
    new_task = Task(description=task_description)
    with tracer.start_as_current_span("adding_task"):
        db.session.add(new_task)
        db.session.commit()
        task_counter.add(1)
        request_counter.add(1, {"endpoint": "add"})
    return redirect(url_for('home'))
```

How does Elastic fit into this?

The Elastic Stack natively supports the OpenTelemetry protocol (OTLP). This means trace data and metrics collected from your applications and infrastructure can be sent directly to the Elastic Stack.

OpenTelemetry API/SDK with Elastic APM agents

Available in Java, .NET, Node.js, and Python



Sending to a collector on Docker

A collector is a component that collects, processes, and exports telemetry data like metrics, logs, and traces, facilitating observability across different systems and platforms.

```
# /tmp/otel-collector-config.yaml
receivers:
  otlp:
    protocols:
      grpc:
        endpoint: 0.0.0.0:4317
      http:
        endpoint: 0.0.0.0:4318

processors:
  batch:
  memory_limiter:
    check_interval: 1s
    limit_mib: 2000

exporters:
  debug:
    verbosity: detailed
  otlp/elastic:
    # Elastic APM server https endpoint without the "https://" prefix
    endpoint: "${env:ELASTIC_APM_SERVER_ENDPOINT}"
    headers:
      # Elastic APM Server secret token
      Authorization: "Bearer ${env:ELASTIC_APM_SECRET_TOKEN}"

service:
  pipelines:
    traces:
      receivers: [otlp]
      exporters: [debug, otlp/elastic]
    metrics:
      receivers: [otlp]
      exporters: [debug, otlp/elastic]
    logs:
      receivers: [otlp]
      exporters: [debug, otlp/elastic]
```

You can configure your collector
via a YAML file





```
pip install opentelemetry-exporter-otlp
```



```
docker run -p 4317:4317 \
-v $(pwd)/otelcollector-config.yaml:/etc/otel-collector-config.yaml \
otel/opentelemetry-collector:latest \
--config=/etc/otel-collector-config.yaml
```





```
export OTEL_PYTHON_LOGGING_AUTO_INSTRUMENTATION_ENABLED=true  
opentelemetry-instrument --service_name otel-collector --logs_exporter otlp flask run -p 8080
```



The Elastic APM OpenTelemetry bridge allows you to create Elastic APM Transactions and Spans, using the OpenTelemetry API.

State of adoption

WORK in PROGRESS

<https://opentelemetry.io/status/>

Closing thoughts

OpenTelemetry (OTel) is an open-source,
vendor-neutral observability framework
designed to integrate with any backend
system.

It's very easy to start instrumenting your code.

It is highly configurable and extensible

It scales well from large applications to
small applications

Next steps



Resources

<https://github.com/JessicaGarson/Introduction-to-OpenTelemetry-with-Flask>



**Let me know if this talk inspires you to build
anything. I'm @JessicaGarson on most
platforms.**

Thank you!